Hacking Machine Learning Systems

Cédric Simal

Motivation

The current state of computer security

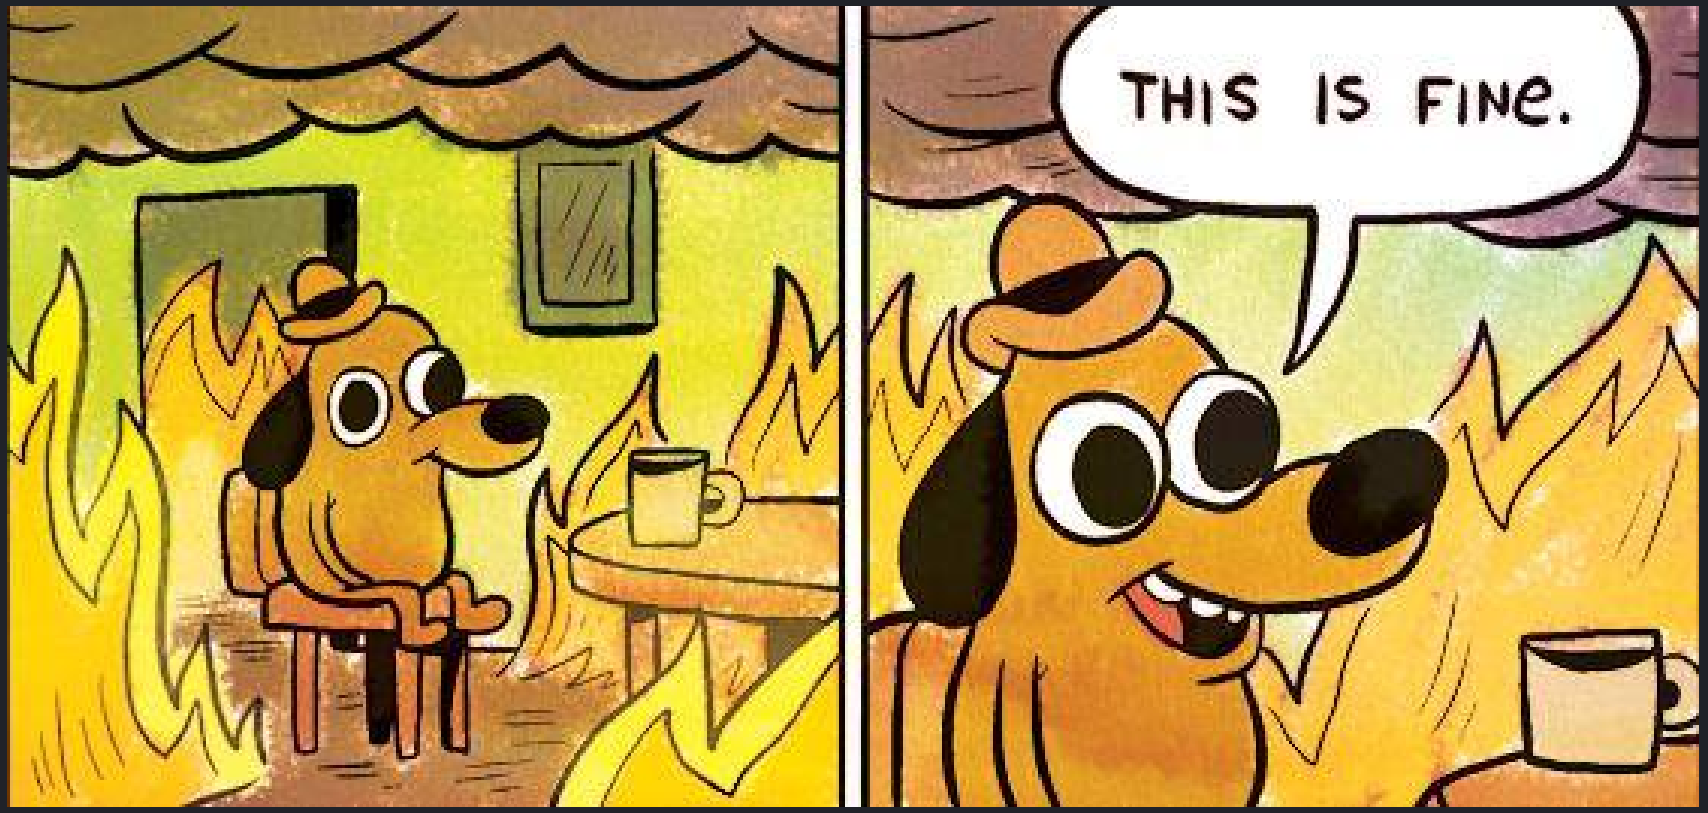

http://gunshowcomic.com/648

Plan

0. Why you should care
1. Security Framework
2. Common Attack types
3. Defenses

All software can be hacked

https://knowyourmeme.com/memes/sites/tay-ai

Why attack ML?

- Degrade model performance
- Find misclassified inputs
- Steal model parameters
- Steal training data

Framework for attacks on ML

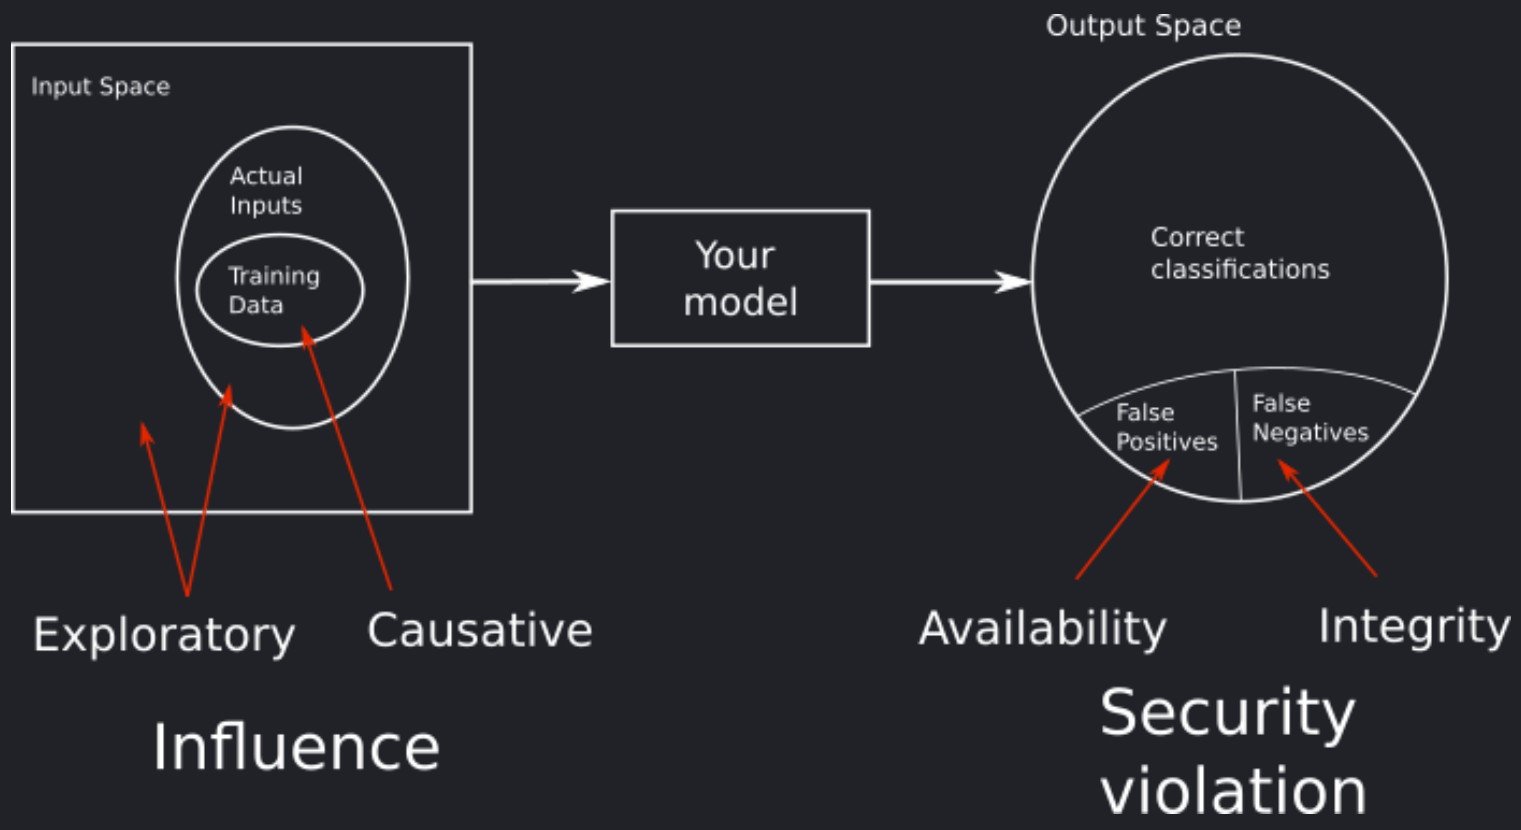

[Barreno et al.]

Data Poisoning

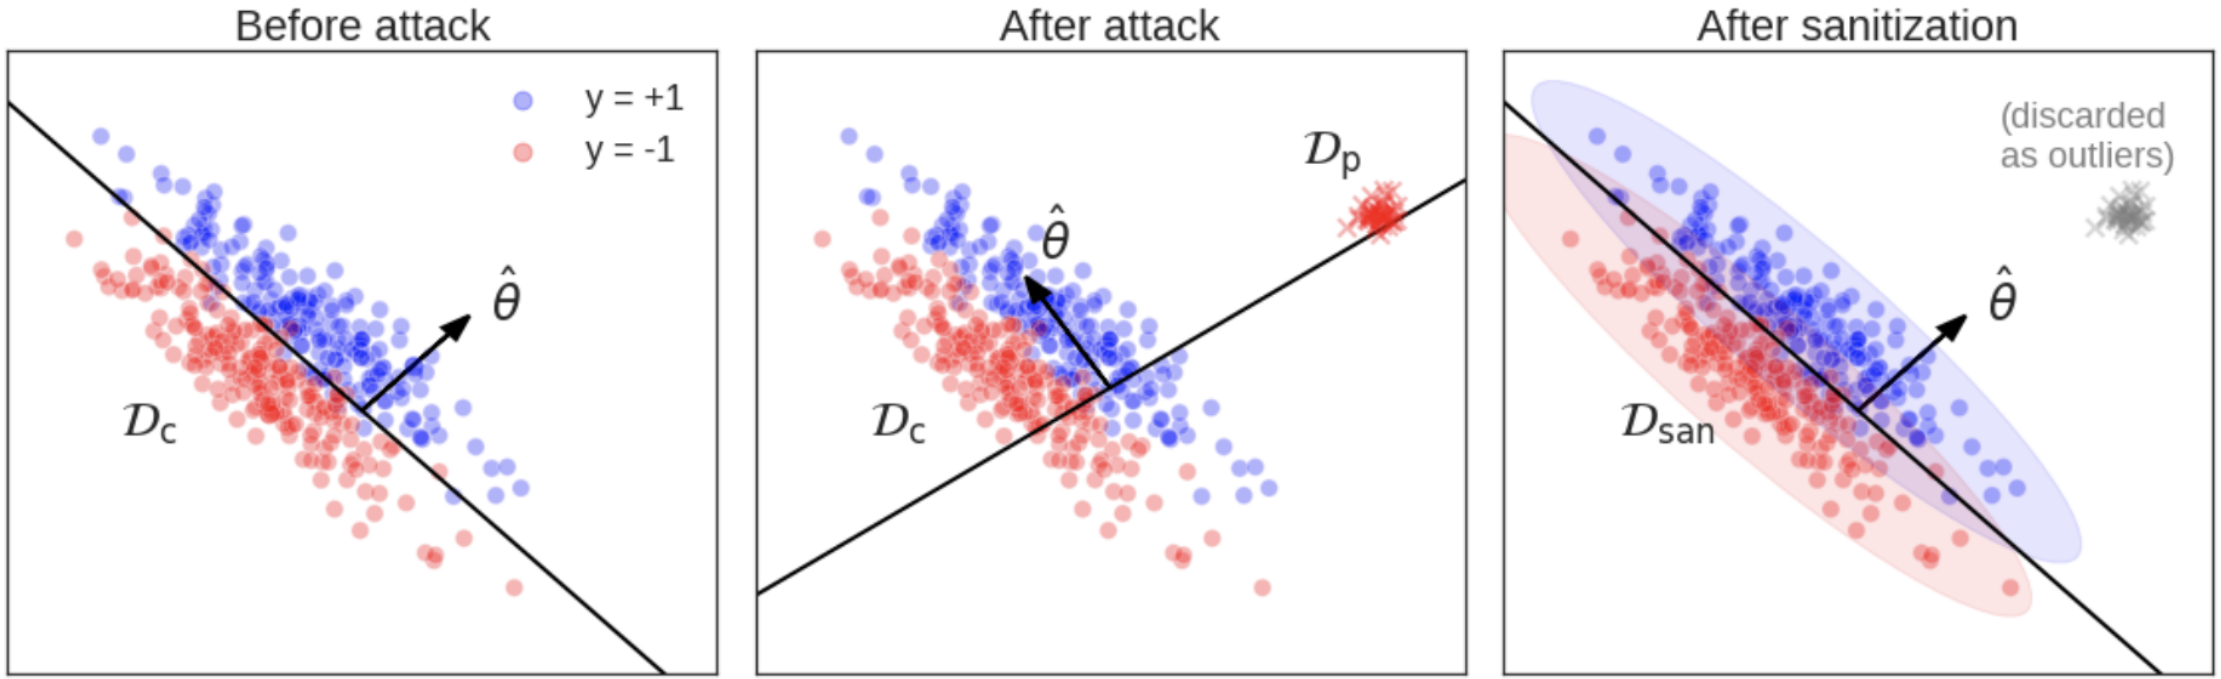

[Koh et al.]

Adversarial Examples

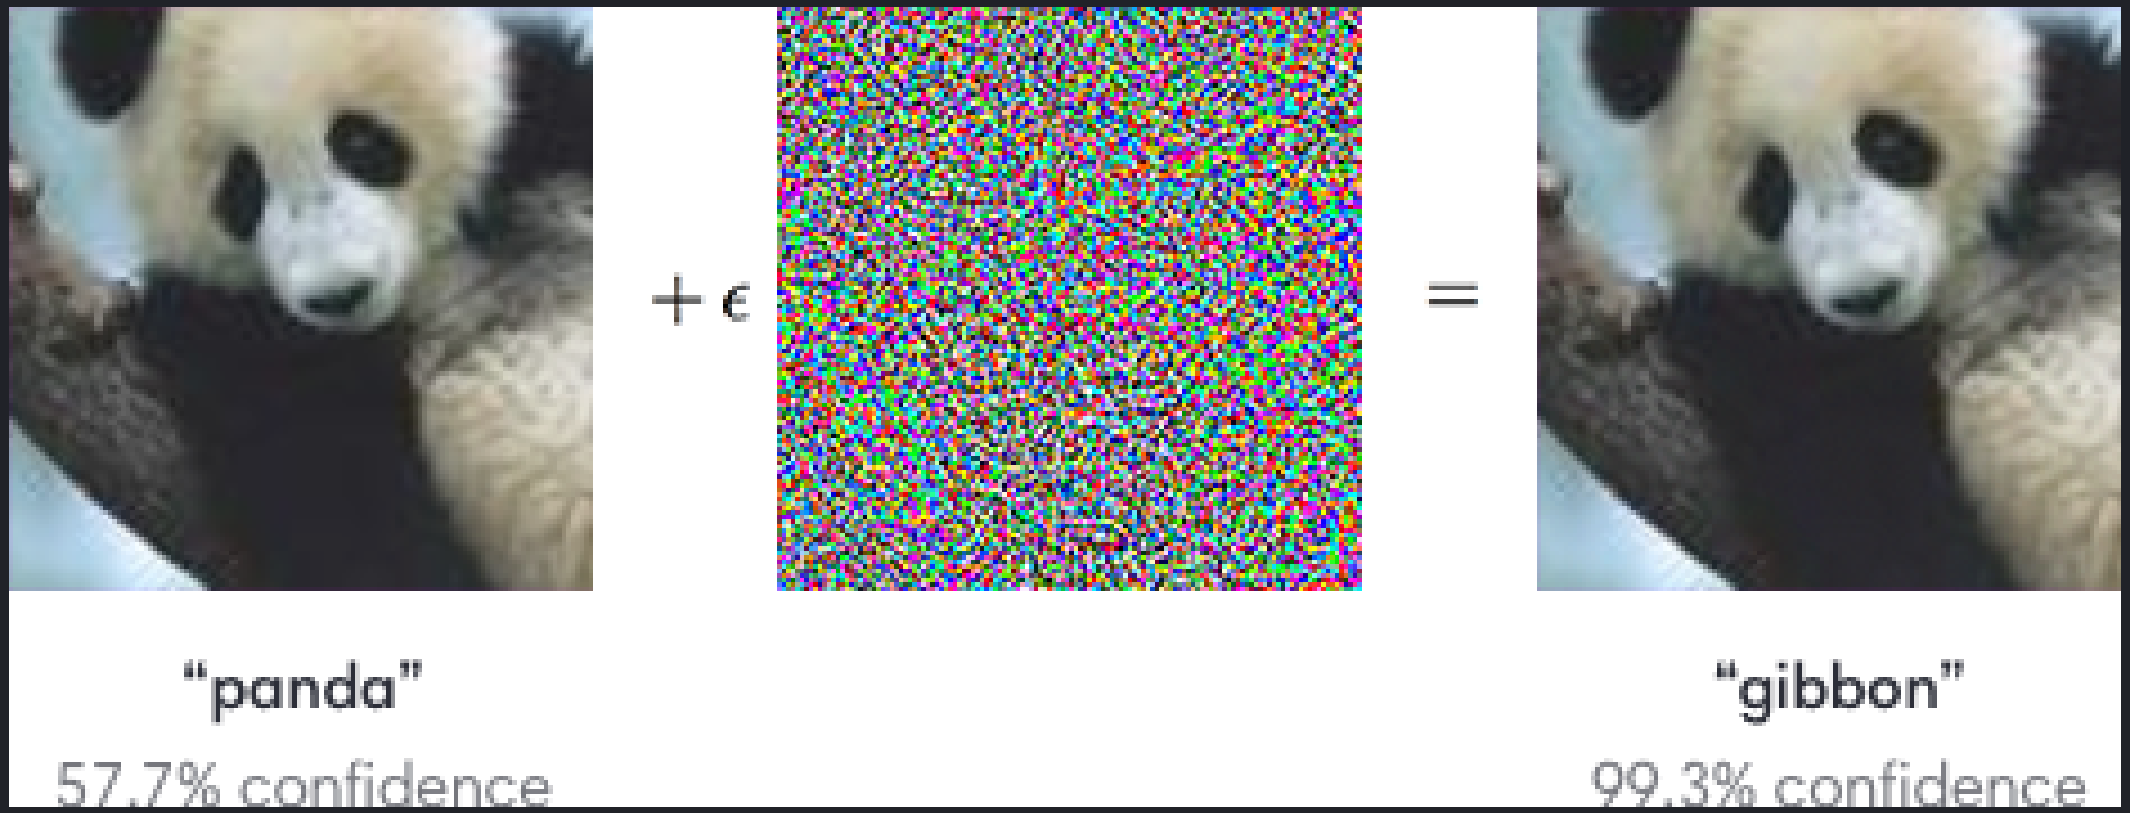

openai.com

Finding Adversarial Examples

- Whitebox : Fast Gradient Sign Method [Goodfellow et al.]
- Blackbox: Surrogate model [Papernot et al.]

Attack Transferability

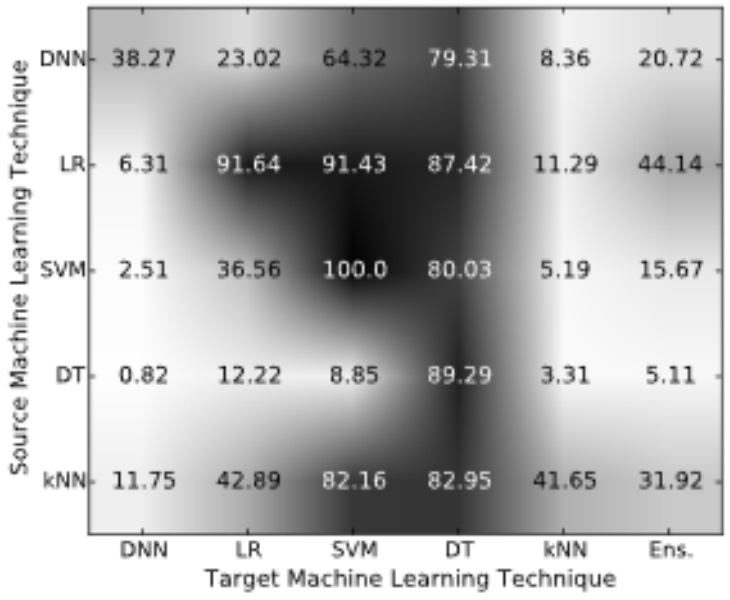

Figure 3: cross-technique Transferability matrix: cell (i, j) is the percentage of adversarial samples crafted to mislead a classifier learned using machine learning technique i that are misclassified by a classifier trained with technique j.

[Papernot et al.]

Differential Privacy

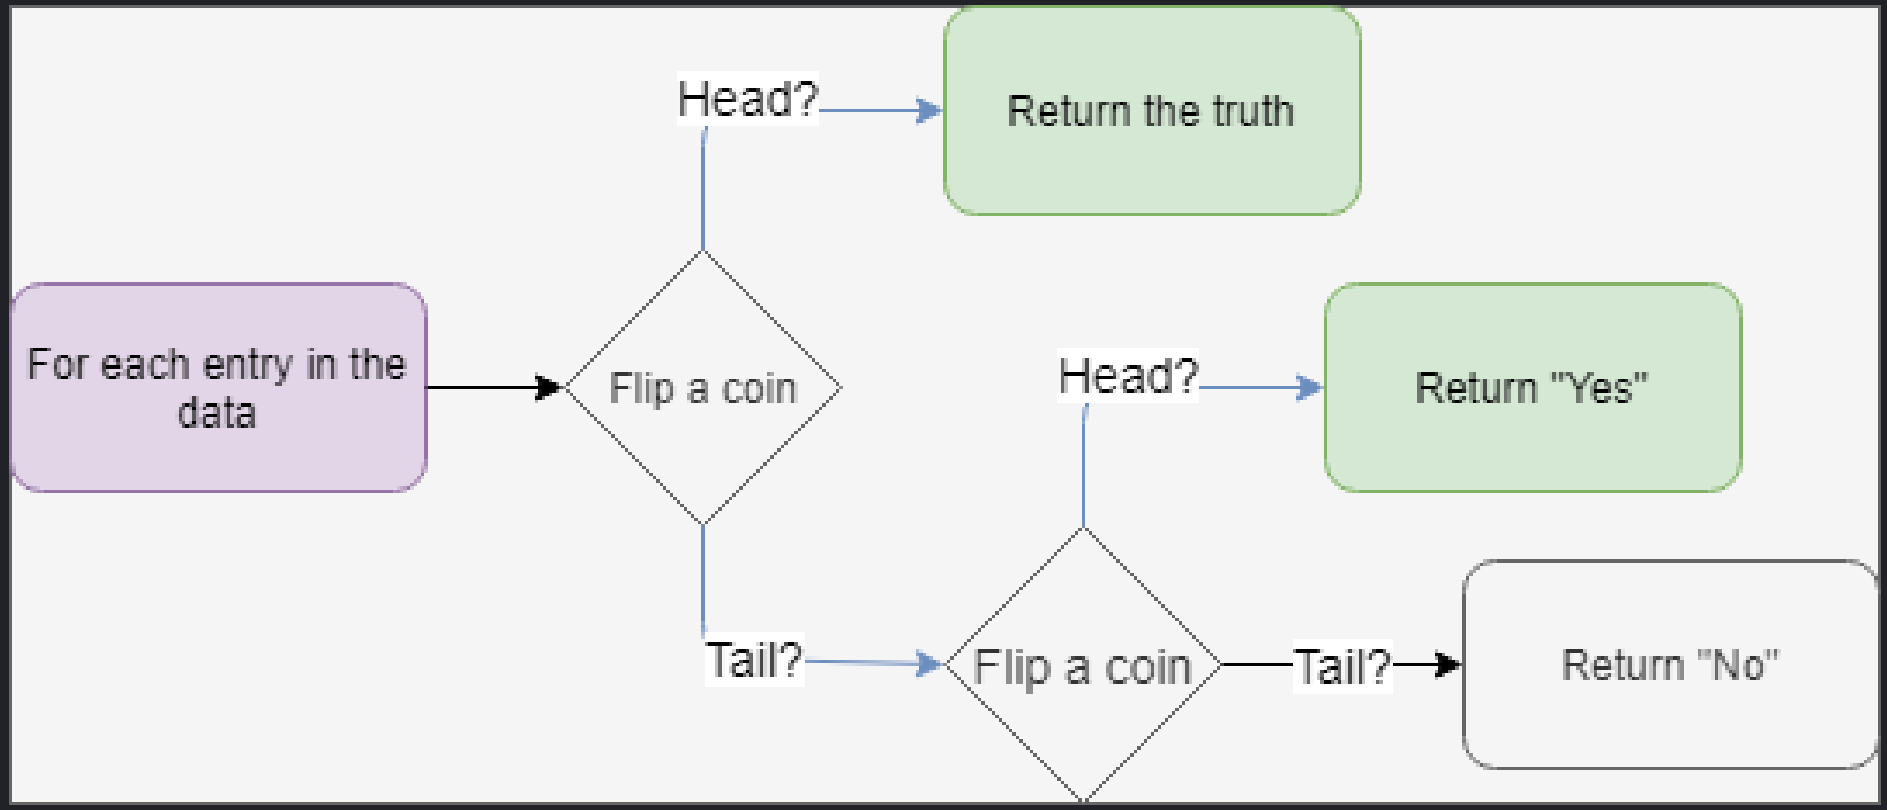

towardsdatascience.com

Defense Strategies

- Reject On Negative Impact (RONI)
- Adversarial Training
- Ensemble methods

Take home message

- Review your data
- Check your inputs
- Include Security at the design stage

Learn more!

Slides and sources at
https://github.com/csimal/ML-Hacking-Presentation